

Representing Data Elements

Physical Layout of Data

In relational terms:

- Field = sequence of bytes representing the value of an attribute in a tuple.
- Record = sequence of bytes divided into fields, representing a tuple.
- File = collection of blocks used to hold a relation = set of tuples or records, respectively.

In object-oriented terms:

- Field represents an attribute or relationship.
- Record represents an object.
- File represents extent of a class.

WS 24/25
Frankfurt UAS

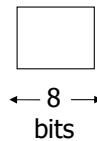
Prof. Dr. Justus Klingemann

The Mapping Problem

What are the data items we want to store?

- a salary
- a name
- a date
- a picture

→ What we have available: Bytes



Implementation of DBMS

Numbers

Integer: 2/4 bytes

e.g., 35 is

00000000

00100011

- Real, floating point
 n bits for mantissa, m for exponent....

WS 24/25
Frankfurt UAS

Prof. Dr. Justus Klingemann

WS 24/25
Frankfurt UAS

Prof. Dr. Justus Klingemann

Characters

→ various coding schemes suggested,
most popular is ascii

Example:

A: 1000001
a: 1100001
5: 0110101
LF: 0001010

Date and Time

Dates

- e.g.: - Integer, # days since Jan 1, 1900
- 8 characters, YYYYMMDD
(not YYMMDD!)
- 7 characters, YYYYDDD
- SQL: YYYY-MM-DD

Time

- e.g. - Integer, seconds since midnight
- characters, HHMMSSFF

Boolean Values

Boolean

e.g., TRUE
FALSE

1	1	1	1	1	1	1
---	---	---	---	---	---	---

0	0	0	0	0	0	0
---	---	---	---	---	---	---

Strings of Characters

- Null terminated
e.g.,

c	a	t	\0	X	X
---	---	---	----	---	---

- Length given
e.g.,

3	c	a	t	X	X	X
---	---	---	---	---	---	---

- Fixed length

Records

Consider fixed-length records first

Record consists of

- Space for each field of the record.
 - Sometimes, it is required to align fields starting at a multiple of 4 or 8.

Example: Employee record

- (1) E#, 2 byte integer
- (2) E_name, 10 char.
- (3) Dept, 2 byte code

55	s m i t h	02
83	j o n e s	01

Schema

Records

Prof. Dr. Justus Klingemann

Record Header

Usually the fields of the record are preceded by a header

Header = space for information about the record, e.g.,

- record format (pointer to schema),
- record length,
- timestamp.

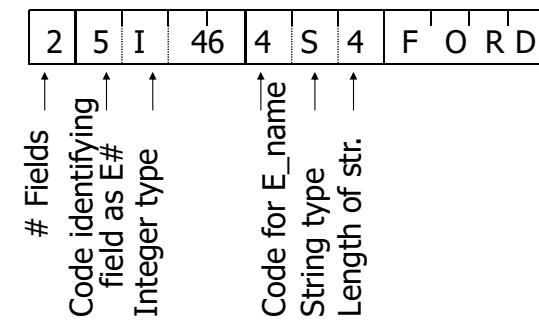
Variable-Length Records

Can occur in case of

- Fields that vary in length
- Repeating fields, e.g., a set of pointers represent a manymany relationship
- Variableformat records: field names are arbitrary
 - Important for selfdescribing data, information integration.

Prof. Dr. Justus Klingemann

Example: variable format and length



Field name codes could also be strings, i.e. TAGS

Example: Repeating Fields

Employee has one or more children

3	E_name: Fred	Child: Sally	Child: Tom
---	--------------	--------------	------------

Internal Organization of Record

a)	*	F1	*	F2	12	F3
----	---	----	---	----	----	----

b)	3	32	5	15	20	F1	F2	F3
	0	1	2	3	4	5	15	20

offsets

Strategy a) Each field is preceded by a number providing its length

Strategy b) The Record header contains a set of pointers to the variable length fields

Hybrid Format

Hybrid format

- one part is fixed, other variable

E.g.: All employees have E#, name, dept
other fields vary.

25	Smith	Toy	2	Hobby:chess	state:retired
# of var fields					

Alternative realization

- Split Records Into Fixed/Variable Parts
- Fixed part has a pointer to space where current value of variable fields can be found.

Placing Records into Blocks

Structure of Blocks:

1. Block header = space for info such as:

- Links to other blocks of a data structure.
- Role info for this block, e.g., for which relation does the block hold tuples?
- Directory of records in the block.
- Block ID.
- Timestamp.

2. Some number of records

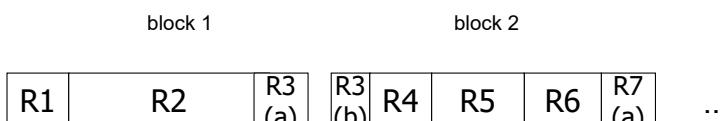
Options for storing records in blocks

- (1) separating records
- (2) spanned vs. unspanned
- (3) mixed record types – clustering
- (4) split records
- (5) sequencing
- (6) addressing

Unspanned: records must be within one block



A spanned record can be divided between two blocks



Separating records

When does a record ends and the next starts?

Block



(a) no need to separate - fixed size recs.

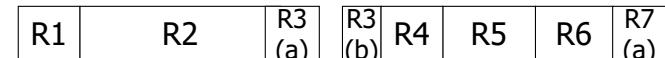
(b) special marker

(c) give record lengths (or offsets)

- within each record
- in block header

Spanned records

Implementation of DBMS



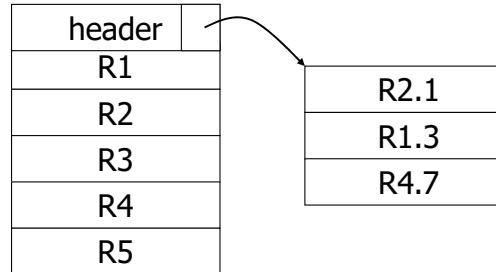
need indication
of partial record
+ “pointer” to rest

need indication
of continuation

Sequencing Options (2)

(c) Overflow area

Records
in sequence



Addressing

How does one refer to records?

Two approaches:

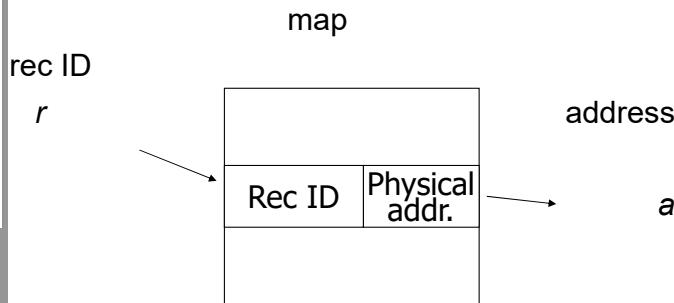
- Physical: sequence of bytes describing location

For example in case of a HDD: device ID, cylinder #, surface #, block # within track, offset within block (for records).

- Logical (indirect): a map table associates abstract ID's, perhaps fixed-length character strings, with physical addresses.

Fully Indirect

E.g., Record ID is arbitrary bit string



Addressing (Cont.)

The map table is itself a relation; Fast access is important.

Physical addresses are more efficient.

Logical addresses allow flexibility:

- records can move or be deleted without dangling pointers.

Common compromise: physical to block level, table of record offsets within blocks.

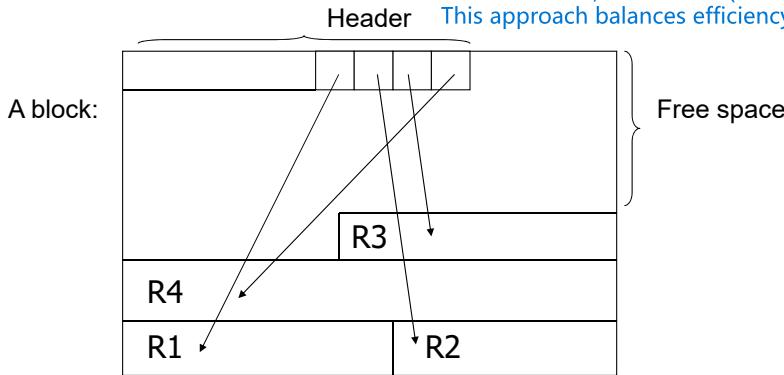
- Tuple-Identifier (TID)
- movement within block: references remain unchanged; only table is modified
- movement to different block: original block contains another TID instead of the record
- regular reorganizations necessary

Implementation of DBMS

Indirection in Block

- Common Compromise in DBMS

DBMSs usually mix both approaches:
Physical addressing to the block level.
Each record is identified by its block number
Logical addressing inside the block.
Within a block, a slot table (offset table) tells
This approach balances efficiency and flexibility.

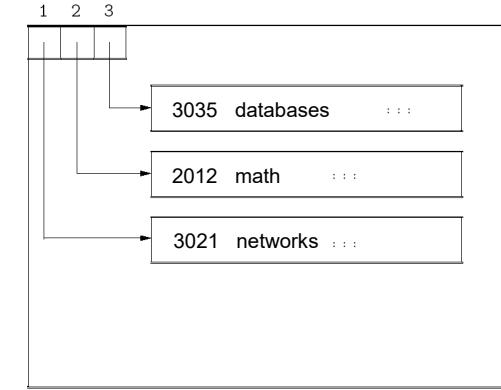
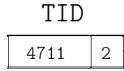


Prof. Dr. Justus Klingemann

WS 24/25
Frankfurt UAS

TID Example

(fast access).
where each record s



Each record is given a TID = (BlockID, SlotID).

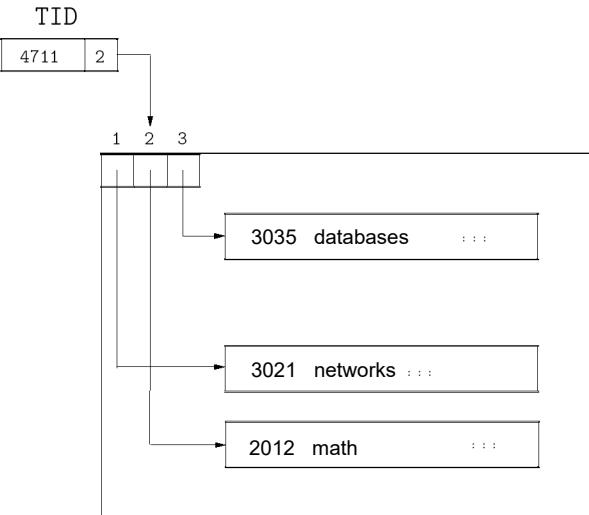
Example: TID = (123, 5) means "record in Block 123, slot 5".

Prof. Dr. Justus Klingemann

page 4711

Implementation of DBMS

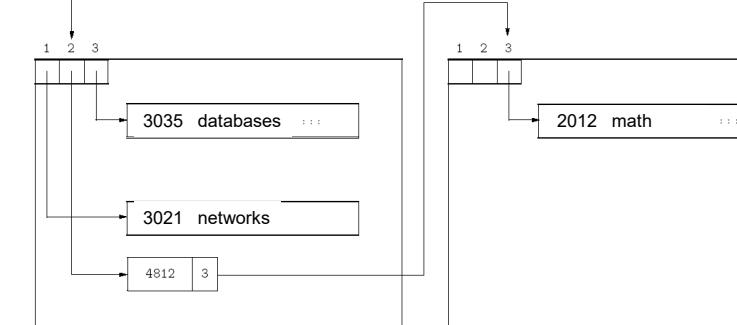
Movement Within Block



Prof. Dr. Justus Klingemann

WS 24/25
Frankfurt UAS

Movement to Different Block



page 4711

page 4812

Prof. Dr. Justus Klingemann

WS 24/25
Frankfurt UAS

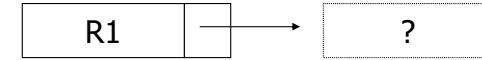
Options for Deletion

- (a) Immediately reclaim space
- (b) Mark deleted
 - Need a way to mark:
 - special characters
 - in map
 - May need chain of deleted records (for re-use)

As usual many tradeoffs

- How expensive is it to move valid record to free space for immediate reclaim?
- How much space is wasted?
 - e.g., deleted records, delete fields, free space chains,...

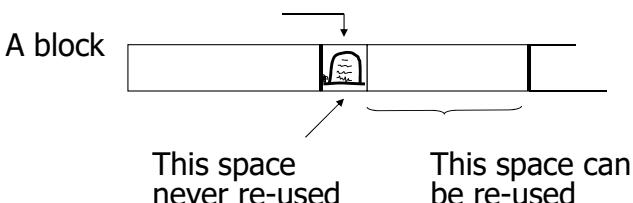
Dangling Pointers



Solution: Tombstones (1)

E.g., leave “MARK” in map or old location

- Physical IDs



Solution: Tombstones (2)

E.g., Leave “MARK” in map or old location

- Logical IDs

map	
ID	LOC
7788	▀

Never reuse
ID 7788 nor
space in map...

Pointer Swizzling

Typical DB structure:

- Data maintained by DBMS, using physical or logical addresses of perhaps 8 bytes.
- Application programs are clients with their own (conventional, virtual-memory) address spaces.

When blocks and records are copied to client's memory, DB addresses can be swizzled = translated to virtualmemory addresses.

- Allows conventional pointer following.
- Especially important in OODBMS, where pointers refer to other objects.

Returning Blocks to Disk

Pointers in the returned block must be unswizzled.

Locate swizzled pointers to block (list has to be managed) and unswizzle.

Swizzling Options

1. Never swizzle. Keep a translation table of DB pointers to local pointers; consult table to follow any DB pointer.

- Problem: time to follow pointers.

2. Automatic swizzling. When a block is copied to memory, replace all its DB pointers by local pointers.

- Problem: large investment if not too many pointerfollowings occur.

3. Swizzle on demand. When a block is copied to memory, do not translate pointers within the block. If we follow a pointer, translate it the first time.

- Problem: requires a bit in pointer fields for DB/local, extra decision at each pointer following.