

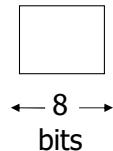
Representing Data Elements

The Mapping Problem

What are the data items we want to store?

- a salary
- a name
- a date
- a picture

→ What we have available: Bytes



Physical Layout of Data

In relational terms:

- Field = sequence of bytes representing the value of an attribute in a tuple.
- Record = sequence of bytes divided into fields, representing a tuple.
- File = collection of blocks used to hold a relation = set of tuples or records, respectively.

In object-oriented terms:

- Field represents an attribute or relationship.
- Record represents an object.
- File represents extent of a class.

Numbers

Integer: 2/4 bytes

e.g., 35 is

00000000

00100011

- Real, floating point
 n bits for mantissa, m for exponent....

Characters

→ various coding schemes suggested,
most popular is ascii

Example:

A: 1000001
a: 1100001
5: 0110101
LF: 0001010

Date and Time

Dates

- e.g.: - Integer, # days since Jan 1, 1900
- 8 characters, YYYYMMDD
(not YYMMDD!)
- 7 characters, YYYYDDD
- SQL: YYYY-MM-DD

Time

- e.g. - Integer, seconds since midnight
- characters, HHMMSSFF

Boolean Values

Boolean

e.g., TRUE
FALSE

1111 1111

0000 0000

Application specific Enumerations

e.g., RED → 1 GREEN → 3
BLUE → 2 YELLOW → 4 ...

➡ Can we use less than 1 byte/code?

Yes, but only if desperate...

Strings of Characters

- Null terminated
e.g.,

c	a	t	\0	X
---	---	---	----	---

- Length given
e.g.,

3	c	a	t	X	X
---	---	---	---	---	---

- Fixed length

Records

Consider fixed-length records first

Record consists of

- Space for each field of the record.
 - Sometimes, it is required to align fields starting at a multiple of 4 or 8.

Example: Employee record

- (1) E#, 2 byte integer
- (2) E_name, 10 char.
- (3) Dept, 2 byte code

Schema

55	s m i t h	02
83	j o n e s	01

Records

Prof. Dr. Justus Klingemann

Record Header

Usually the fields of the record are preceded by a header

Header = space for information about the record, e.g.,

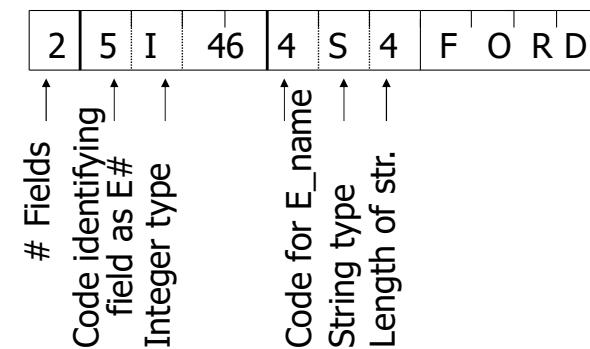
- record format (pointer to schema),
- record length,
- timestamp.

Variable-Length Records

Can occur in case of

- Fields that vary in length
- Repeating fields, e.g., a set of pointers represent a manymany relationship
- Variableformat records: field names are arbitrary
 - Important for selfdescribing data, information integration.

Example: variable format and length



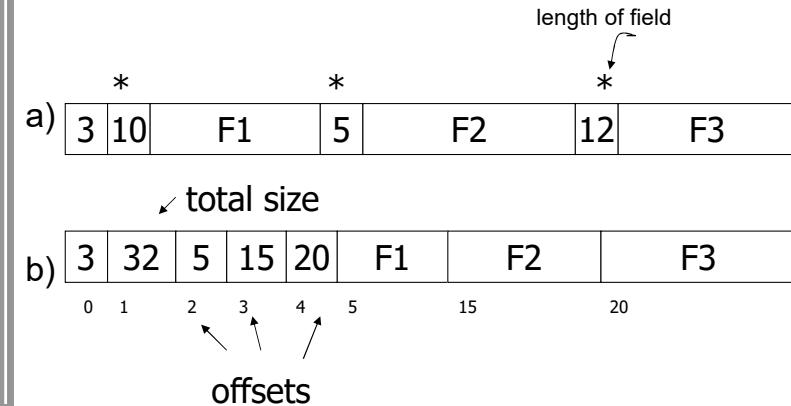
Field name codes could also be strings, i.e. TAGS

Example: Repeating Fields

Employee has one or more children

3	E_name: Fred	Child: Sally	Child: Tom
---	--------------	--------------	------------

Internal Organization of Record



Strategy a) Each field is preceded by a number providing its length

Strategy b) The Record header contains a set of pointers to the variable length fields

Hybrid Format

Hybrid format

- one part is fixed, other variable

E.g.: All employees have E#, name, dept
other fields vary.

25	Smith	Toy	2	Hobby:chess	state:retired
----	-------	-----	---	-------------	---------------

of var fields

Alternative realization

- Split Records Into Fixed/Variable Parts
- Fixed part has a pointer to space where current value of variable fields can be found.

Placing Records into Blocks

Structure of Blocks:

- Block header = space for info such as:
 - Links to other blocks of a data structure.
 - Role info for this block, e.g., for which relation does the block hold tuples?
 - Directory of records in the block.
 - Block ID.
 - Timestamp.
- Some number of records

Options for storing records in blocks

- (1) separating records
 - (2) spanned vs. unspanned
 - (3) mixed record types – clustering
 - (4) split records
 - (5) sequencing
 - (6) addressing

WS 24/25
Frankfurt UAS

Prof. Dr. Justus Klingema

Separating records

When does a record ends and the next starts?

Bloc|



- (a) no need to separate - fixed size recs
 - (b) special marker
 - (c) give record lengths (or offsets)
 - within each record
 - in block header

WS 24/25
Frankfurt

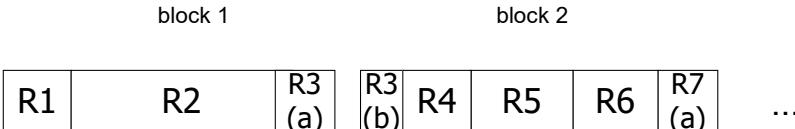
Prof. Dr. Justus Klingemann

Spanned vs. Unspanned

Unspanned: records must be within one block



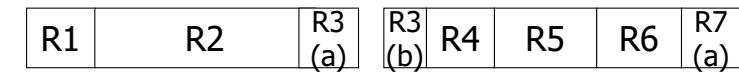
A spanned record can be divided between two blocks



WS 24/25
Frankfurt UAS

Prof. Dr. Justus Klingema

Spanned records



need indication
of partial record
+ “pointer” to res

need indication
of continuation

卷之三

WS 24/25
Frankfurt

Prof. Dr. Justus Klingemann

Spanned vs. Unspanned

Unspanned is much simpler, but may waste space...

Spanned essential if

$$\text{record size} > \text{block size}$$

Mixed record types

Mixed: records of different types (e.g. EMPLOYEE, DEPT) allowed in same block

e.g., a block

EMP	e1	DEPT	d1	DEPT	d2	
-----	----	------	----	------	----	--

Why do we want to mix?

Answer: CLUSTERING

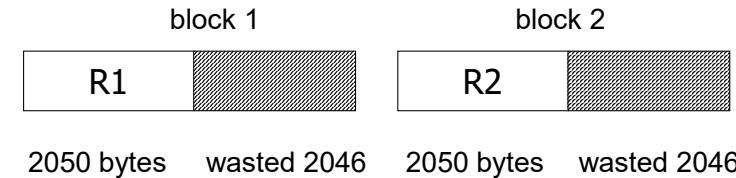
Records that are frequently accessed together should be in the same block

Example

10^6 records

each of size 2,050 bytes (fixed)

block size = 4096 bytes



Total wasted $\approx 2 \times 10^9$ Utilization $\approx 50\%$

Total space $\approx 4 \times 10^9$

Example

Q1: select A#, C_NAME, C_CITY, ...
from DEPOSIT, CUSTOMER
where DEPOSIT.C_NAME =
CUSTOMER.C_NAME

a block

CUSTOMER,NAME=SMITH
DEPOSIT,NAME=SMITH
DEPOSIT,NAME=SMITH

Options for storing records in blocks

If Q1 frequent, clustering is good

But if Q2 frequent

```
Q2: SELECT *\n      FROM CUSTOMER
```

CLUSTERING IS COUNTER PRODUCTIVE

Split records

Typically for hybrid format

Fixed part in one block

Variable part in another block