

3.9 Exercises for Section 3

Exercise 3.1: Show what happens to the buckets in Fig. 20 if the following insertions and deletions occur:

- i. Records g through j are inserted into buckets 0 through 3, respectively.
- ii. Records a and b are deleted.
- iii. Records k through n are inserted into buckets 0 through 3, respectively.
- iv. Records c and d are deleted.

Exercise 3.2: We did not discuss how deletions can be carried out in a linear or extensible hash table. The mechanics of locating the record(s) to be deleted should be obvious. What method would you suggest for executing the deletion? In particular, what are the advantages and disadvantages of restructuring the table if its smaller size after deletion allows for compression of certain blocks?

! Exercise 3.3: The material of this section assumes that search keys are unique. However, only small modifications are needed to allow the techniques to work for search keys with duplicates. Describe the necessary changes to insertion, deletion, and lookup algorithms, and suggest the major problems that arise when there are duplicates in each of the following kinds of hash tables: (a) simple (b) linear (c) extensible.

INDEX STRUCTURES

! Exercise 3.4: Some hash functions do not work as well as theoretically possible. Suppose that we use the hash function on integer keys i defined by $h(i) = i^2 \bmod B$, where B is the number of buckets.

- a) What is wrong with this hash function if $B = 10$?
- b) How good is this hash function if $B = 16$?
- c) Are there values of B for which this hash function is useful?

Exercise 3.5: In an extensible hash table with n records per block, what is the probability that an overflowing block will have to be handled recursively; i.e., all members of the block will go into the same one of the two blocks created in the split?

Exercise 3.6: Suppose keys are hashed to four-bit sequences, as in our examples of extensible and linear hashing in this section. However, also suppose that blocks can hold three records, rather than the two-record blocks of our examples. If we start with a hash table with two empty blocks (corresponding to 0 and 1), show the organization after we insert records with hashed keys:

- a) 0000, 0001, ..., 1111, and the method of hashing is extensible hashing.
- b) 0000, 0001, ..., 1111, and the method of hashing is linear hashing with a capacity threshold of 100%.
- c) 1111, 1110, ..., 0000, and the method of hashing is extensible hashing.
- d) 1111, 1110, ..., 0000, and the method of hashing is linear hashing with a capacity threshold of 75%.

Exercise 3.7: Suppose we use a linear or extensible hashing scheme, but there are pointers to records from outside. These pointers prevent us from moving records between blocks, as is sometimes required by these hashing methods. Suggest several ways that we could modify the structure to allow pointers from outside.

!! Exercise 3.8: A linear-hashing scheme with blocks that hold k records uses a threshold constant c , such that the current number of buckets n and the current number of records r are related by $r = ckn$. For instance, in Example 24 we used $k = 2$ and $c = 0.85$, so there were 1.7 records per bucket; i.e., $r = 1.7n$.

- a) Suppose for convenience that each key occurs exactly its expected number of times.⁷ As a function of c , k , and n , how many blocks, including overflow blocks, are needed for the structure?

⁷This assumption does not mean all buckets have the same number of records, because some buckets represent twice as many keys as others.

INDEX STRUCTURES

- b) Keys will not generally distribute equally, but rather the number of records with a given key (or suffix of a key) will be *Poisson distributed*. That is, if λ is the expected number of records with a given key suffix, then the actual number of such records will be i with probability $e^{-\lambda} \lambda^i / i!$. Under this assumption, calculate the expected number of blocks used, as a function of c , k , and n .

! Exercise 3.9: Suppose we have a file of 1,000,000 records that we want to hash into a table with 1000 buckets. 100 records will fit in a block, and we wish to keep blocks as full as possible, but not allow two buckets to share a block. What are the minimum and maximum number of blocks that we could need to store this hash table?