



### Question 1a:

Consider the following queries on relations R(A,B) and S(B,C):

- **Query 1:** SELECT \* FROM R, S WHERE R.B = S.B AND R.A = 5
- **Query 2:** SELECT \* FROM R, S WHERE R.B = S.B AND R.A = 5 AND S.C = 10

The following two plans are being considered:



Relation R has **50,000 tuples** with **20 tuples per disk block**,  $V(R,A)=5$ , and  $V(R,B)=10$ .

Similarly, S has **25,000 tuples** with **25 tuples per disk block**,  $V(S,B)=4$ , and  $V(S,C)=15$ . Assume values are distributed over possible  $V(\text{Relation}, \text{Attribute})$  values (not over possible domain values).

In this exercise, we make the following additional assumptions:

- The join is implemented as a **hash-join**;
- The intermediate result produced by operation (1.1) is **not written to disk**;
- Hash buckets are stored on disk;
- The final result is kept in **main-memory**;
- There is enough memory to execute the hash join algorithm.

a) How many disk I/O's does **Query 1** require?


b) Also assume that the result of neither (2.1) nor (2.2) is stored to disk. How many disk I/O's does **Query 2** require?

### Question 1b:

Consider the following queries on relations R(A,B) and S(B,C):

- **Query 1:** SELECT \* FROM R, S WHERE R.B = S.B AND R.A = 3
- **Query 2:** SELECT \* FROM R, S WHERE R.B = S.B AND R.A = 3 AND S.C = 7

The following two plans are being considered:



Relation R has **40,000 tuples** with **15 tuples per disk block**,  $V(R,A)=4$ , and  $V(R,B)=8$ .

Similarly, SS has **20,000 tuples** with **20 tuples per disk block**,  $V(S,B)=3$ , and  $V(S,C)=10$ . Assume values are distributed over possible  $V(\text{Relation}, \text{Attribute})$  values (not over possible domain values).

In this exercise, we make the following additional assumptions:

- The join is implemented as a **hash-join**;
- The intermediate result produced by operation (1.1) is **not written to disk**;
- Hash buckets are stored on disk;
- The final result is kept in **main-memory**;
- There is enough memory to execute the hash join algorithm.

a) How many disk I/O's does **Query 1** require?


b) Also assume that the result of neither (2.1) nor (2.2) is stored to disk. How many disk I/O's does **Query 2** require?

### Question 1c:

Consider the following queries on relations R(A,B) and S(B,C):

- **Query 1:** SELECT \* FROM R, S WHERE R.B = S.B AND R.A = 9
- **Query 2:** SELECT \* FROM R, S WHERE R.B = S.B AND R.A = 9 AND S.C = 12

The following two plans are being considered:



Relation RR has **70,000 tuples** with **25 tuples per disk block**,  $V(R,A)=7$ , and  $V(R,B)=14$ .

Similarly, SS has **35,000 tuples** with **35 tuples per disk block**,  $V(S,B)=6$ , and  $V(S,C)=25$ . Assume values are distributed over possible  $V(\text{Relation}, \text{Attribute})$  values (not over possible domain values).

In this exercise, we make the following additional assumptions:

- The join is implemented as a **hash-join**;
- The intermediate result produced by operation (1.1) is **not written to disk**;
- Hash buckets are stored on disk;
- The final result is kept in **main-memory**;
- There is enough memory to execute the hash join algorithm.

a) How many disk I/O's does **Query 1** require?

b) Also assume that the result of neither (2.1) nor (2.2) is stored to disk. How many disk I/O's does **Query 2** require?

---

### Key Features of These Questions:

1. **Query Plans:**
  - Each query involves a join between two relations RR and SS, with additional filtering conditions.
2. **Relation Statistics:**
  - Each relation has a specified number of tuples, tuples per block, and distinct values for attributes.
3. **Assumptions:**
  - The join is implemented as a hash-join.
  - Intermediate results are not written to disk.
  - Hash buckets are stored on disk.
  - The final result is kept in main-memory.

#### 4. Disk I/O Calculations:

- The goal is to compute the number of disk I/O's required for each query under the given assumptions.
- 

**2a)** Suppose  $B(R) = 30000$ ,  $B(S) = 70000$ , and the number of main memory blocks is  $M = 150$ . We want to perform the **natural join of R and S using a hash join algorithm**.

- a) How many partitions are needed to ensure that each partition of R fits in memory?
  - b) Describe the process of building and probing the hash table and compute the number of I/O operations required.
  - c) How many main memory blocks would be required to perform a one-pass hash join?
- 

**2b)** Suppose  $B(R) = 10000$ ,  $B(S) = 25000$ , and we have  $M = 200$  main memory blocks available. We wish to join R and S using a **sort-merge join algorithm**, assuming both relations are initially unsorted.

- a) How many sorting passes are required for each relation?
  - b) Describe how the sorted relations are merged and the number of I/O's required.
  - c) What is the minimum number of main memory blocks needed for an efficient sort-merge join?
- 

**2c)** Suppose  $B(R) = 15000$ ,  $B(S) = 40000$ , and the available **main memory blocks** are  $M = 120$ . We want to perform a **nested-loop join of R and S**.

- a) How many block accesses are required for a **tuple nested-loop join**?
  - b) How many block accesses are required for a **block nested-loop join**?
  - c) What is the minimum number of memory blocks needed to perform a **single-pass block nested-loop join**?
- 

#### Example 1: External Merge Sort for Join Preparation

Suppose we have two relations, **R and S**, with the following properties:  $B(R)=15000B$ ,  $B(S)=40000$ , The number of available **main memory blocks** is  $M=80$ .

- Both relations are **clustered but unsorted**, and we want to sort them before applying a **merge join**.

#### Questions:

- a) How many passes do we need to sort R and S separately using external merge sort?
  - b) Describe the execution of the merge join and calculate the I/O cost.
  - c) How many main memory blocks would be needed to perform a one-pass join?
- 

#### Example 2: Sort-Merge Join with Different Memory Constraints

Consider two relations:

- $B(R)=30000$ ,  $B(S)=70000$ , and  $M = 120$ .

- Both relations are **not sorted** and must be sorted before using a **merge join**.
- Assume **R and S are clustered** on disk.

#### Questions:

- Compute the number of passes needed to sort **R and S** using **external merge sort**.
  - How many total I/O operations are required for the full sort-merge join process?
  - If we increase the available memory to  $M=250M = 250M=250$ , how would this affect the sorting and join efficiency?
- 

#### Example 3: Comparing Sort-Merge and Block Nested Loop Joins

We have two relations:  $B(R)=50000$  and  $B(S)=90000$ , The number of **main memory blocks** available is  $M=60$ .

- The relations are **not sorted**, and we want to perform a join.

#### Questions:

- How many passes are required if we first sort both relations and use a **merge join**?
  - Compare the total **I/O cost** of using a **sort-merge join** versus a **block nested loop join** under these constraints.
  - What is the minimum number of memory blocks needed for a **one-pass join**, and is it feasible in this case?
- 

#### Question 1: Sort-Merge Join with Clustered Relations

Suppose  $B(R)=15000$ ,  $B(S)=30000$ , and the number of main memory blocks is  $M=201$ . We want to perform the natural join of **R** and **S** using a **sort-merge join algorithm**. Both relations are clustered but neither is sorted.

- How many passes are required to sort **R** and **S**?
  - Describe how the join is executed and calculate the total number of I/O operations required.
  - How many main memory blocks would be needed to perform a **one-pass join**?
- 

#### Question 2: Hash Join with Unclustered Relations

Suppose  $B(R)=25000$ ,  $B(S)=40000$ , and the number of main memory blocks is  $M=151$ . We want to perform the natural join of **R** and **S** using a **hash join algorithm**. Both relations are unclustered.

- How many passes are required for the hash join?
  - Describe how the join is executed and calculate the total number of I/O operations required.
  - How many main memory blocks would be needed to perform a **one-pass join**?
- 

#### Question 3: Block-Nested Loop Join with Sorted Relations

Suppose  $B(R)=10000$ ,  $B(S)=50000$ , and the number of main memory blocks is  $M=51$ . We want to perform the natural join of **R** and **S** using a **block-nested loop join algorithm**. Relation **R** is sorted, but **S** is not sorted.

- How many passes are required for the block-nested loop join?

- b) Describe how the join is executed and calculate the total number of I/O operations required.  
c) How many main memory blocks would be needed to perform a **one-pass join**?
- 

**3a)** We want to sort the tuples of a relation **S** on a given key. The following information is known:

- **S contains 200,000 tuples**, i.e.,  $T(S) = 200000$ .
- The **size of a block** on disk is **4096 bytes**.
- The **size of each tuple** in **S** is **512 bytes**.
- Relation **S** is **clustered**, meaning each disk block is fully occupied with **S** tuples.
- The **size of the sort key** is **40 bytes**.
- A **record pointer** is **8 bytes**.

Answer the following questions:

- If we use a **two-pass sorting algorithm**, what is the minimum amount of **main memory** (in number of blocks) required?
  - What is the **I/O cost** of the two-pass sorting algorithm, including the cost of writing the sorted file to disk?
  - If instead of sorting entire tuples, we sort only the **<key, recordPointer>** pairs and merge them later using record pointers to retrieve tuples, what is the **I/O cost** in this case?
- 

**3b)** Consider a relation **T** that needs to be sorted on a particular key. Given the following information:

- **T contains 150,000 tuples**, i.e.,  $T(T) = 150000$ .
- The **disk block size** is **4096 bytes**.
- Each tuple in **T** has a size of **256 bytes**.
- Relation **T** is **clustered**.
- The **sort key** is **48 bytes**.
- A **record pointer** is **8 bytes**.

Answer the following questions:

- How many **blocks** are required to store relation **T**?
  - What is the **cost of two-pass sorting** for this relation in terms of **disk I/Os**?
  - If we use a **three-pass sorting algorithm** instead of two-pass sorting, what is the **minimum main memory required** in terms of blocks?
- 

**3c)** Suppose we need to sort a **relation Q** based on a given key. The relation has the following properties:

- **Q contains 250,000 tuples**, i.e.,  $T(Q) = 250000$ .
- The **block size on disk** is **4096 bytes**.
- The **size of each tuple** in **Q** is **320 bytes**.
- The relation is **clustered**.
- The **sort key** is **36 bytes**.
- A **record pointer** is **8 bytes**.

Answer the following:

- How many **disk blocks** are required to store **Q**?
- If a **two-pass sorting algorithm** is used, what is the **minimum memory required** (in terms of blocks)?

c) What is the **disk I/O cost** if we sort only **<key, recordPointer>** pairs first, then retrieve tuples during the merge phase?

---

## Question 1: Merge Join with Limited Memory

Consider two relations A and B with the following properties:

- $B(A) = 15,000$  blocks,  $B(B) = 40,000$  blocks, and the number of available main memory blocks  $M = 50$ .
- Both relations are **clustered but unsorted**.

### Questions:

- How many passes are required to sort A and B before performing a merge join?
- Describe the steps involved in performing the **merge join** and determine the **total I/O cost**.
- How much **main memory** would be needed to perform the join in **one pass** instead of multiple passes?

### Concept Explanation:

- Merge join requires that both relations be **sorted first** if they are not already sorted.
- External merge sort requires **log-based passes** over the data to sort it.
- **Sorting cost** is  $2B \times (\text{number of passes})$  since each block is read and written multiple times.
- Once sorted, the **merge join** can be performed in **one pass**, requiring only  $B(A) + B(B)$  I/Os.
- If we had **enough memory to fit one entire relation**, the join could be completed in **one pass**.

## Question 2: Hash Join with Limited Partitions

Consider relations X(P, Q) and Y(Q, R):

- $B(X) = 25,000$  blocks,  $B(Y) = 60,000$  blocks.
- Available main memory blocks  $M = 200$ .
- We need to perform a **hash join** between X and Y.

### Questions:

- How many **partitions** will be created during the **partitioning phase** of the hash join?
- What is the **total I/O cost** of performing the **hash join**?
- If  $M = 5,000$ , would hash join still be the best option, or would another join method be preferable?

### Concept Explanation:

- Hash join consists of two phases:
  1. **Partitioning:** Each relation is divided into **M-1 partitions**, ensuring that each partition fits in memory.
  2. **Probing:** Each partition of X is matched with its corresponding partition in Y.
- **Partitioning cost:** Each relation is **read and written once**.
- **Probing cost:** Each partition is **read again** and joined.
- **Total I/O cost** for hash join:  
$$\downarrow$$
$$3(B(X) + B(Y))$$

• If M is large enough to fit one relation entirely in memory, a **one-pass join** is possible.

## Question 3: Comparing Nested-Loop Join and Block Nested-Loop Join

Consider two relations M(S, T) and N(T, U):

- $B(M) = 35,000$  blocks,  $B(N) = 80,000$  blocks.
- Available main memory blocks  $M = 400$ .

### Questions:

- a) What is the I/O cost if we use a simple nested-loop join?
- b) How does using a block nested-loop join improve performance, and what is its I/O cost?
- c) If  $M$  were increased to 10,000, would block nested-loop join still be the best choice?

### Concept Explanation:

- Simple Nested-Loop Join (NLJ):

- For each block of  $M$ , scan  $N$  entirely.
- I/O Cost:

$$B(M) \times B(N) + B(M)$$

- Block Nested-Loop Join (BNLJ):

- Load  $M-1$  blocks of  $M$  into memory at once.
- Scan  $N$  using fewer passes.
- I/O Cost:

$$\left( \frac{B(M)}{M-1} \right) \times B(N) + B(M)$$

- Larger  $M$  reduces cost, making BNLJ more efficient than simple nested-loop join.

---

Here are three similar questions focusing on **external sorting**, **I/O optimization**, and **efficient sorting techniques**:

---

### Question 3a: Sorting a Relation with a Limited Memory Budget

We need to sort the tuples of a relation  $S$  on a given key. The following details are provided:

- Relation  $S$  contains 200,000 tuples, i.e.,  $T(S) = 200,000$ .
- The size of a block on disk is 8,000 bytes.
- The size of each tuple is 500 bytes.
- Relation  $S$  is clustered.
- The sort key size is 40 bytes, and a record pointer is 8 bytes.

### Questions:

- a) What is the **minimum number of main memory blocks** required to sort the relation using a **two-pass sorting algorithm**?
- b) What is the **total disk I/O cost** for sorting **S** using the two-pass algorithm, including writing the final sorted relation back to disk?
- c) If we modify the sorting algorithm to sort **only** the **<key, recordPointer>** pairs instead of the full tuples (using a technique similar to an index sort), how does this affect the **I/O cost**?

#### Concept Explanation:

- The **two-pass sorting algorithm** requires enough memory to handle **B(S) / M** sorted runs in one merge pass.
  - **Sorting only key-pointer pairs** reduces **I/O costs** since fewer data blocks are read/written.
  - The **total I/O cost** consists of reading, writing, and merging the sorted runs.
- 

#### Question 3b: External Merge Sort with Large Relation

We need to **sort relation T** using **external merge sort**. Given:

- **Relation T has 500,000 tuples**, i.e.,  $T(T) = 500,000$ .
- **Block size is 4,096 bytes**.
- Each **tuple is 512 bytes**.
- **T is not clustered**, meaning tuples may not be packed efficiently in blocks.
- The **available main memory has 500 blocks**.

#### Questions:

- a) How many **passes** are required for sorting **T** using **external merge sort**?
- b) What is the **total I/O cost**, including the merge phase?
- c) If we store the **sorted relation as an index** instead of full tuples, how much **storage** is saved, and what is the **new I/O cost**?

#### Concept Explanation:

- The **number of passes** in external merge sort depends on the **initial sorted runs** and how many runs we can merge at once.
  - The **total I/O cost** is computed by summing all **read/write operations** in each phase.
  - Using **index-based sorting** (sorting keys instead of tuples) reduces **storage and I/O costs**.
- 

#### Question 3c: Optimizing External Sorting with Indexed Sorting

We consider a scenario where we optimize sorting by only sorting **keys and pointers**.

- **Relation U has 1,000,000 tuples**, i.e.,  $T(U) = 1,000,000$ .
- **Block size = 8,192 bytes**.
- Each **tuple is 1,024 bytes**.
- **Sort key = 64 bytes**, and each **record pointer is 8 bytes**.

- The relation is clustered.
- Available main memory has 1,000 blocks.

#### Questions:

- How many **sorted runs** are produced in the **initial pass** of external sorting?
- If we only sort **<key, recordPointer>** pairs, what is the **new disk I/O cost**?
- How does **sorting only index entries** instead of full tuples improve performance in DBMS query processing?

#### Concept Explanation:

- Sorting **only key-pointer pairs** significantly reduces the **amount of data written to disk**.
- The **merge process** then retrieves **tuples from disk only when needed**, reducing **random I/O operations**.
- This technique is often used in **B+ tree index sorting** for efficient **query lookups**.