


Question 1: Query Optimization with Hash Join

Consider the following queries on relations $R(A, B)$ and $S(B, C)$:

Query 1: `SELECT * FROM R, S WHERE R.B = S.B AND R.A = 5`

Query 2: `SELECT * FROM R, S WHERE R.B = S.B AND R.A = 5 AND S.C = 10`

The following two plans are being considered:



Relation R has 50,000 tuples with 20 tuples per block, $V(R, A) = 10$ and $V(R, B) = 15$.

Similarly, S has 40,000 tuples with 25 tuples per block, $V(S, B) = 8$ and $V(S, C) = 25$. Assume values are distributed over possible $V(\text{Relation}, \text{Attribute})$ values (not over possible domain values).

In this exercise, we make the following additional assumptions:

- The join is implemented as a hash-join;
- The intermediate result produced by operation (1.1) is not written to disk;
- Hash buckets are stored on disk;
- The final result is kept in main-memory;
- There is enough memory to execute the hash join algorithm.

a) How many I/O's does Query 1 require?

b) Also assume that the result of neither (2.1) nor (2.2) is stored to disk. How many I/O's does Query 2 require?

Q.1

$$B(R) = \frac{25,000}{16 \text{ bR}} = 2500 \text{ bR}$$

$$B(S) = \frac{16,000}{8 \text{ bR}} = 1600 \text{ bR}$$

$$T(1,1) = T(2,1) = \frac{T(R)}{\text{W.A. A7}_R} = \frac{50,000}{16} = 5000 \text{ tuples}$$

$$T(2,2) = \frac{T(S)}{V(S,C)} = \frac{40,000}{8} = 1600 \text{ tuples}$$

$$B(1,1) = B(2,1) = \frac{250}{25 \text{ fp/bR}} = 250 \text{ bR}$$

$$B(2,2) = \frac{1600 \text{ fp}}{25 \text{ fp/bR}} = 64 \text{ bR}$$

a) query-1: we need 2500 I/O to read R and execute selection.

we need 250 I/O to write the hash to disk and again 250 I/O's to read from disk in second phase.

From S, we need 1600 I/O to read S and create hash function.

We need another 1600 to make hash to hash buckets to disk and again 1600 I/O's to read from disk in second phase of the hash join:

total: [7800 I/O's]

b) Query-2: we already have 3000 I/O from left branch. we need 1600 I/O to read S and execute selection. we need 64 to write hash to disk and again 64 to read from disk:

in total: [4728 I/O's]


Question 1b: Query Optimization with Hash Join

Consider the following queries on relations $R(A, B)$ and $S(B, C)$:

Query 1: SELECT * FROM R, S WHERE R.B = S.B AND R.A = 10

Query 2: SELECT * FROM R, S WHERE R.B = S.B AND R.A = 10 AND S.C = 20

The following two plans are being considered:



Relation R has 100,000 tuples with 25 tuples per block, $V(R, A) = 12$ and $V(R, B) = 25$.

Similarly, S has 80,000 tuples with 30 tuples per block, $V(S, B) = 15$ and $V(S, C) = 40$.

Assume values are distributed over possible $V(\text{Relation}, \text{Attribute})$ values (not over possible domain values).

In this exercise, we make the following additional assumptions:

- The join is implemented as a hash-join;
- The intermediate result produced by operation (1.1) is not written to disk;
- Hash buckets are stored on disk;
- The final result is kept in main-memory;
- There is enough memory to execute the hash join algorithm.

a) How many I/O's does Query 1 require?

b) Also assume that the result of neither (2.1) nor (2.2) is stored to disk. How many I/O's does Query 2 require?


Question 1c: Query Optimization with Different Selectivity

Consider the following queries on relations $R(A, B)$ and $S(B, C)$:

Query 1: SELECT * FROM R, S WHERE R.B = S.B AND R.A = 3

Query 2: SELECT * FROM R, S WHERE R.B = S.B AND R.A = 3 AND S.C = 15

The following two plans are being considered:



Relation R has 70,000 tuples with 15 tuples per block, $V(R, A) = 7$ and $V(R, B) = 20$. Similarly, S has 50,000 tuples with 20 tuples per block, $V(S, B) = 10$ and $V(S, C) = 30$. Assume values are distributed over possible $V(\text{Relation}, \text{Attribute})$ values (not over possible domain values).

In this exercise, we make the following additional assumptions:

- The join is implemented as a hash-join;
- The intermediate result produced by operation (1.1) is not written to disk;
- Hash buckets are stored on disk;
- The final result is kept in main-memory;
- There is enough memory to execute the hash join algorithm.

a) How many I/O's does Query 1 require?

b) Also assume that the result of neither (2.1) nor (2.2) is stored to disk. How many I/O's does Query 2 require?

Exercise 1a

Consider the following queries on relations **T(X, Y)** and **U(Y, Z)**:

- **Query 1:** `SELECT * FROM T, U WHERE T.Y = U.Y AND T.X = 5`
- **Query 2:** `SELECT * FROM T, U WHERE T.Y = U.Y AND T.X = 5 AND U.Z = 10`

The following two plans are being considered:

Query Plan 1

1. **Filter:** Apply selection **T.X = 5** on **T**
2. **Join:** Perform a hash join with **U** on **T.Y = U.Y**

Query Plan 2

1. **Filter:** Apply selection **T.X = 5** on **T**
2. **Filter:** Apply selection **U.Z = 10** on **U**
3. **Join:** Perform a hash join on **T.Y = U.Y**

Assumptions:

- **T** has **80,000** tuples, **U** has **40,000** tuples
- **T** has **20 tuples per block**, **U** has **40 tuples per block**
- **V(T, Y) = 10**, **V(U, Y) = 8**, **V(U, Z) = 25**
- The join is implemented as a **hash join**
- Intermediate results are **not written to disk**
- There is **enough memory** to execute the join

Questions:

- a) How many I/O's does Query 1 require?
 - b) How many I/O's does Query 2 require?
-

Exercise 1b

Consider the following queries on relations **A(M, N)** and **B(N, P)**:

- **Query 1:** `SELECT * FROM A, B WHERE A.N = B.N AND A.M = 3`
- **Query 2:** `SELECT * FROM A, B WHERE A.N = B.N AND A.M = 3 AND B.P = 7`

Query Plan 1

1. **Filter:** Apply selection **A.M = 3** on **A**
2. **Join:** Perform a hash join with **B** on **A.N = B.N**

Query Plan 2

1. **Filter:** Apply selection **A.M = 3** on **A**
2. **Filter:** Apply selection **B.P = 7** on **B**
3. **Join:** Perform a hash join on **A.N = B.N**

Assumptions:

- **A** has **120,000** tuples, **B** has **60,000** tuples
- **A** has **15 tuples per block**, **B** has **30 tuples per block**

- $V(A, N) = 12$, $V(B, N) = 10$, $V(B, P) = 30$
- The join is implemented as a **hash join**
- Intermediate results are **not written to disk**
- There is **enough memory** to execute the join

Questions:

- How many I/O's does Query 1 require?
 - How many I/O's does Query 2 require?
-

Exercise 1c)

Consider the following queries on relations **X(A, B)** and **Y(B, C)**:

- **Query 1:** `SELECT * FROM X, Y WHERE X.B = Y.B AND X.A = 2`
- **Query 2:** `SELECT * FROM X, Y WHERE X.B = Y.B AND X.A = 2 AND Y.C = 4`

Query Plan 1

1. **Filter:** Apply selection **X.A = 2** on **X**
2. **Join:** Perform a hash join with **Y** on **X.B = Y.B**

Query Plan 2

1. **Filter:** Apply selection **X.A = 2** on **X**
2. **Filter:** Apply selection **Y.C = 4** on **Y**
3. **Join:** Perform a hash join on **X.B = Y.B**

Assumptions:

- **X** has **90,000** tuples, **Y** has **50,000** tuples
- **X** has **25 tuples per block**, **Y** has **35 tuples per block**
- $V(X, B) = 15$, $V(Y, B) = 12$, $V(Y, C) = 28$
- The join is implemented as a **hash join**
- Intermediate results are **not written to disk**
- There is **enough memory** to execute the join

Questions:

- How many I/O's does Query 1 require?
- How many I/O's does Query 2 require?

Question 2a)

Suppose we have two relations **A** and **B** with the following properties:

- **$B(A) = 30,000$, $B(B) = 60,000$** , The number of available **main memory blocks** is $M = 150$
- Both relations are **clustered but unsorted**

We want to perform a **merge join** on **A** and **B** using the available memory.

Questions:

- How many sorting passes are required before the merge phase?
 - Describe how the merge join is executed and compute the total I/O cost.
-

Question 2b)

Consider two relations **T(X, Y)** and **U(Y, Z)**:

- **$B(T) = 25,000$, $B(U) = 75,000$, $M = 120$** main memory blocks are available
- Both relations are **clustered but not sorted**

We aim to join **T** and **U** on **T.Y = U.Y** using the **merge join algorithm**.

Questions:

- Determine the number of passes required to sort each relation before merging.
 - Explain the execution of the merge join and calculate the total I/O cost.
-

Question 2c)

Given the relations **R(A, B)** and **S(B, C)**: **$B(R) = 50,000$, $B(S) = 90,000$** $M = 200$ main memory blocks

- Both relations are **clustered but initially unsorted**, We need to perform a **merge join** on **R** and **S**.

Questions:

- Compute the number of sorting passes needed before merging.
 - Describe the merge join execution and determine the I/O cost.
-

Question 3a: Join with Duplicate Elimination

Relations $R(a, b)$ and $S(b, c)$ have r and s tuples, respectively. In both relations, attribute b takes values 1, 2, ..., 20 only, with equal probability. We also assume that both relations have duplicates, and on average a tuple that appears at all appears **four times** (i.e., eliminating duplicates reduces the size of R or S by a factor of 4). We wish to compute, without duplicates, those tuples in the natural join of R and S that have $b = 20$. Here are three possible plans for answering this query:

Plan A	Plan B	Plan C
$T1 := \sigma_{b=20}(R)$	$T4 := \sigma_{b=20}(R)$	$T7 := \delta(R)$
$T2 := \delta(S)$	$T5 := \sigma_{b=20}(S)$	$T8 := \sigma_{b=20}(S)$
$T3 := T1 \text{ JOIN } T2$	$T6 := T4 \text{ JOIN } T5$	$T9 := T7 \text{ JOIN } T8$
$\text{Ans} := \delta(T3)$	$\text{Ans} := \delta(T6)$	$\text{Ans} := \delta(T9)$

Provide for each plan a formula for the expected sum of sizes of the intermediate relations depending on the values of r and s .

Question 3b: Join with Specific Attribute Value

Relations $R(a, b)$ and $S(b, c)$ have r and s tuples, respectively. In both relations, attribute b takes values 1, 2, ..., 15 only, with equal probability. We also assume that both relations have duplicates, and on average a tuple that appears at all appears **five times** (i.e., eliminating duplicates reduces the size of R or S by a factor of 5). We wish to compute, without duplicates, those tuples in the natural join of R and S that have $b = 15$. Here are three possible plans for answering this query:

Plan A	Plan B	Plan C
$T1 := \sigma_{b=15}(R)$	$T4 := \sigma_{b=15}(R)$	$T7 := \delta(R)$
$T2 := \delta(S)$	$T5 := \sigma_{b=15}(S)$	$T8 := \sigma_{b=15}(S)$
$T3 := T1 \text{ JOIN } T2$	$T6 := T4 \text{ JOIN } T5$	$T9 := T7 \text{ JOIN } T8$
$\text{Ans} := \delta(T3)$	$\text{Ans} := \delta(T6)$	$\text{Ans} := \delta(T9)$

Provide for each plan a formula for the expected sum of sizes of the intermediate relations depending on the values of r and s .

Question 3c: Join with Different Duplicate Factors

Relations $R(a, b)$ and $S(b, c)$ have r and s tuples, respectively. In both relations, attribute b takes values $1, 2, \dots, 25$ only, with equal probability. We also assume that both relations have duplicates, and on average a tuple that appears at all appears **two times** in R and **three times** in S (i.e., eliminating duplicates reduces the size of R by a factor of 2 and S by a factor of 3). We wish to compute, without duplicates, those tuples in the natural join of R and S that have $b = 25$. Here are three possible plans for answering this query:

Plan A	Plan B	Plan C
$T1 := \sigma_{b=25}(R)$	$T4 := \sigma_{b=25}(R)$	$T7 := \delta(R)$
$T2 := \delta(S)$	$T5 := \sigma_{b=25}(S)$	$T8 := \sigma_{b=25}(S)$
$T3 := T1 \text{ JOIN } T2$	$T6 := T4 \text{ JOIN } T5$	$T9 := T7 \text{ JOIN } T8$
Ans := $\delta(T3)$	Ans := $\delta(T6)$	Ans := $\delta(T9)$

Provide for each plan a formula for the expected sum of sizes of the intermediate relations depending on the values of r and s .

Explanation of Concepts in the Questions

1. **Natural Join:** Combines tuples from R and S based on a common attribute b .
2. **Duplicate Elimination (δ):** Removes duplicate tuples from a relation.
3. **Selection (σ):** Filters tuples based on a condition (e.g., $b = 10$).
4. **Intermediate Relations:** Temporary relations created during query execution (e.g., $T1, T2, T3$).
5. **Expected Sizes:** The average sizes of intermediate relations, calculated based on the number of tuples (r and s) and the duplicate factors.

Question 3a:

Consider two relations $R(A, B)$ and $S(B, C)$ with sizes $|R| = 50,000$ tuples and $|S| = 25,000$ tuples. The attribute B has a uniform distribution of values. We aim to compute the natural join $R \bowtie S$.

Three different query plans are proposed:

- **Plan X:** First, sort both R and S on B and then perform a merge join.
- **Plan Y:** Use a hash join with a hash table built on the smaller relation S .
- **Plan Z:** Perform a nested loop join with R as the outer relation.

- a) For each plan, estimate the number of I/Os required.
 - b) Which plan is optimal given a memory size of $M = 200$ pages?
-

Question 3b:

Given relations $R(A, B)$ and $S(B, C)$, we want to evaluate the following query efficiently:

```
SELECT * FROM R, S WHERE R.B = S.B AND R.A < 100 AND S.C > 50;
```

We consider three alternative query execution plans:

- **Plan A:** Apply selection predicates on R and S before joining.
- **Plan B:** Perform the join first, then apply selection predicates.
- **Plan C:** Use an index join, assuming an index exists on $S.B$.

- a) What are the expected intermediate result sizes for each plan?
 - b) Which plan minimizes the I/O cost assuming uniform data distribution?
-

Question 3c:

Suppose we have relations $R(A, B)$ with 80,000 tuples and $S(B, C)$ with 40,000 tuples. The number of unique B values in R and S are 4,000 and 2,000, respectively.

We need to compute the join $R \bowtie S$ using different strategies:

- **Plan 1:** Hash join, with S as the build relation.
- **Plan 2:** Sort-merge join.
- **Plan 3:** Block nested loop join, with R as the outer relation.

- a) Compute the estimated number of disk I/Os required for each plan.
 - b) If the available memory is limited to 500 pages, which plan is the most efficient?
-

4a)

Consider two relations R(A, B) and S(B, C) with sizes $B(R) = 30,000$ and $B(S) = 40,000$ blocks, respectively. The available main memory has $M = 150$ blocks. We want to perform a **one-pass join** between R and S.

- a) How many memory blocks are required to perform the join in a single pass?
 - b) If one-pass join is not feasible, what would be the best alternative join method given the memory constraints?
-

4b)

Given two relations R(X, Y) and S(Y, Z) stored in $B(R) = 10,000$ and $B(S) = 20,000$ blocks, we have $M = 250$ main memory blocks available. Both relations are **unsorted and unclustered**.

- a) Can we perform a **one-pass join** with the given memory?
 - b) If not, how many passes would a **two-pass sort-merge join** require, and what would be the estimated I/O cost?
-

Question 4c:

Suppose we have two relations R(P, Q) and S(Q, R) stored in $B(R) = 50,000$ and $B(S) = 80,000$ blocks. The main memory available for the join operation is $M = 300$ blocks.

- a) What is the minimum number of memory blocks required to perform a **one-pass hash join**?
 - b) If we were to use a **partitioned hash join**, how many partitions would be required, and how would they be distributed across memory?
-

Question 5:

2) Consider a clustered relation R(A, B, C, D) that has a clustering index on A and a non-clustering index on each of the other attributes. The relevant parameters are: $B(R) = 1000$, $T(R) = 5000$, $V(R, A) = 20$, $V(R, B) = 1000$, $V(R, C) = 5000$, and $V(R, D) = 500$. Give the best query plan for the following selection and the corresponding number of disk I/O's. You can ignore the cost for accessing the index.
 $\sigma_{A=10 \text{ AND } C=2 \text{ AND } D=3}(R)$

Question 5a:

Consider a clustered relation R(A, B, C, D) that has a clustering index on B and non-clustering indexes on each of the other attributes. The relevant parameters are:

- $B(R) = 5000$, $T(R) = 10000$, $V(R, A) = 50$, $V(R, B) = 200$, $V(R, C) = 1000$, and $V(R, D) = 4000$.

Determine the best query plan and the number of disk I/Os for the following selection:

$$\sigma_{B=5 \text{ AND } C=20}(R)$$

Question 5b:

Given a relation $S(X, Y, Z, W)$ with a clustering index on X and a non-clustering index on Z , the database parameters are:

- $B(S) = 8000$, $T(S) = 40000$, $V(S, X) = 100$, $V(S, Y) = 500$, $V(S, Z) = 2500$, and $V(S, W) = 10000$.

Identify the optimal query execution plan and calculate the number of disk I/Os required to evaluate the query:

$$\sigma_{X=3 \text{ AND } Z=100}(S)$$

Question 5c:

Suppose we have a relation $T(P, Q, R, S, T)$ where:

- A **clustering index** exists on P ,
- A **non-clustering index** exists on T ,
- $B(T) = 12000$, $T(T) = 60000$, $V(T, P) = 150$, $V(T, Q) = 5000$, $V(T, R) = 1000$, $V(T, S) = 2000$, $V(T, T) = 100$.

Estimate the best way to process the following selection query and compute the corresponding number of disk I/Os:

$$\sigma_{P=50 \text{ AND } T=3}(T)$$
