Implementation of DBMS
Exercise Sheet 15
Klingemann, WS 2024 / 2025

1) Consider a clustered relation R(A, B, C, D) that has a clustering index on A and a non-clustering index on each of the other attributes. The relevant parameters are: B(R) = 1000, T(R) = 5000, V(R, A) = 20, V(R, B) = 1000, V(R, C) = 5000, and V(R, D) = 500. Give the best query plan for the following selection and the corresponding number of disk I/O's. You can ignore the cost for accessing the index.

$\sigma_{A=10 \text{ AND } C=2 \text{ AND } D=3}(R)$