

DBMS Tutorial 13.02.2019

Topics

Index Join and Cost of Selections

Index Join

Assumptions : R has an index on the join attribute A. Index is in memory. **R and S are contiguous.**

Algorithm:

1. Read S
 - a. For each tuple in S probe index to find if $R.A = S.A$ exists.
 - i. If it does, fetch it from disk, and Join
 - ii. Each tuple from S will join with $T(R)/V(R,A)$ tuples of R (expected result of a selection).
 - iii. Assuming the index to be non-clustering, $T(R)/V(R,A)$ records can be fetched in as many IOs
 - iv. Assuming the index to be clustering: $T(R)/V(R,A)$ records can be fetched in $B(R)/V(R,A)$ IOs

Cost = Reading S + retrieving tuples of R whenever an index entry exists

$$= B(S) + T(S) \times T(R)/V(R,A) \text{ for non-clustering index}$$

$$= B(S) + T(S) \times B(R)/V(R,A) \text{ for clustering index}$$

Index Join

Assumptions : R has an index on the join attribute A. Index is in memory. **R and S are contiguous.**

Algorithm:

1. Read S
 - a. For each tuple in S probe index to find if $R.A = S.A$ exists.
 - i. If it does, fetch it from disk, and Join
 - ii. Each tuple from S will join with $T(R)/V(R,A)$ tuples of R (expected result of a selection).
 - iii. Assuming the index to be non-clustering, $T(R)/V(R,A)$ records can be fetched in as many IOs
 - iv. Assuming the index to be clustering: $T(R)/V(R,A)$ records can be fetched in $B(R)/V(R,A)$ IOs

Cost = Reading S + retrieving tuples of R whenever an index entry exists

$$= B(S) + T(S) \times T(R)/V(R,A) \text{ for non-clustering index}$$

$$= B(S) + T(S) \times B(R)/V(R,A) \text{ for clustering index}$$

In case of a join, when one of the relations has an index on join attribute (here R), we read the other relation completely(here S) and for each value of join attribute in S, we probe the index on R to find if such a value exists in R. Expected result from R is similar to a selection query on R.

Index Join Example

Assumptions : R has an index on the join attribute A. Index is in memory. S is contiguous.

Examples:

Case 1. R.A is primary key and S.A is foreign key

Case 2. $V(R,A) = 100$ with non-clustering index

Case 3. $V(R,A) = 100$ with clustering index

Remember:

Relations R and S

$T(R) = 60000$

$T(S) = 2000$

Block size = 4000 bytes

$S(R) = 100$ bytes

$S(S) = 500$ bytes

Main Memory = 126 Buffers

Index Join Example

Assumptions : R has an index on the join attribute A. Index is in memory. S is contiguous. Cost in terms of IOs

Examples:

Case 1. R.A is primary key and R.S is foreign key, ie $V(R,A) = T(R)$

Number of possible matches for each tuple of S = $T(R)/V(R,A) = 1$

Cost = $250 + 2000 \times 1$

Case 2. $V(R,A) = 100$ with non-clustering index

Number of possible matches for each S.A = $T(R)/V(R,A) = 60000/100$

Cost = $250 + 2000 \times 600$

Case 3. $V(R,A) = 100$ with clustering index

Number of possible matches for each S.A = $T(R)/V(R,A) = 60000/100$

but now they can be found in $B(R)/V(R,A) = 1500/100$ blocks!

Cost = $250 + 2000 \times 15$

Notes

Probing Index for R.A = '10'

1. Probing Index means checking the index for a given key. For this the index needs to be in memory. If it's not kept in memory probing the index will be an IO cost.
2. Expected number of tuples for such a query is $T(R)/V(R,A)$
 - a. If the index is clustering we can get all the tuples (using the pointers from the index) in $B(R)/V(R,A)$ IOs.
 - b. If the index is non-clustering we can get all the tuples (using the pointers from the index) in $T(R)/V(R,A)$ IOs.

Example from DSCB

Suppose $B(R) = 1000$, and $T(R) = 20,000$. (This means R has 20,000 tuples, packed at most 20 to a block). Let A be one of the attributes of R , suppose there is an index on A , and consider the operation $\sigma_{A=0}(R)$. Here are some possible situations and the worst-case number of disk I/O 's required.

We shall ignore the cost of accessing the index blocks in all cases.

- a) R is clustered but we do not use the index. Cost = 1000 disk I/O 's.
- b) If R is not clustered and we do not use the index. Cost = 20000 disk I/O 's
- c) If $V(R, A) = 100$ and the index is clustering. Cost = $1000/100 = 10$ disk I/O 's, (plus whatever is needed to access the index).
- d) If $V(R,A) = 10$ and the index is non-clustering. Cost = $20,000/10 = 2000$ disk I/O 's.
(Notice that this cost is higher than scanning the entire relation R , if R is clustered but the index is not.)
- e) If $V(R, A) = 20,000$, i.e., A is a key, then the index-based algorithm takes 1 disk I/O (plus whatever is needed to access the index, regardless of whether the index is clustering or not.)

Exercise 1

$T(R) = 10000$, $B(R) = 500$, $V(R,A) = 30$. Calculate

- a) the number of tuples expected in the result, and
- b) the IOs it will take to retrieve the result, for the query $\sigma_{A=10}(R)$ in the following cases :

1. R has a clustering index on A.
2. R has a non-clustering index on A
3. R has no index on A but is a clustered/contiguous relation.
4. R is non-contiguous.

Exercise 1 Ans

$T(R) = 10000$, $B(R) = 500$, $V(R,A) = 30$. Calculate

- a) the number of tuples expected in the result, and
- b) the IOs it will take to retrieve the result, for the query $\sigma_{A=10}(R)$ in the following cases:

1. R has a clustering index on A.
 - a. Tuples: $T(R)/V(R,A) = 10000/30$
 - b. IOs: $B(R)/V(R,A) = 500/30$
2. R has a non-clustering index on A
 - a. Tuples: $T(R)/V(R,A) = 10000/30$
 - b. IOs: $T(R)/V(R,A) = 10000/30$
3. R has no index on A but is a clustered/contiguous relation
 - a. Tuples: $T(R)/V(R,A) = 10000/30$
 - b. IOs: $B(R) = 500$
4. R is non-contiguous.
 - a. Tuples: $T(R)/V(R,A) = 10000/30$
 - b. IOs: $T(R) = 10000$

Exercise 2

$T(R) = 10000$, $B(R) = 400$, $V(R,A) = 50$, $V(R,B)=30$, $V(R,C)=100$. R has a clustering index on A and non-clustering index on B .

Calculate the cost in terms of a) the number of tuples expected in the result and b) the IOs it will take to retrieve the result in the following cases (assume index to be in memory) :

1. $\sigma_{A<10}(R)$
2. $\sigma_{B=20}(R)$
3. $\sigma_{C=40}(R)$

What would be a cost effective plan to execute this query:

4. $\sigma_{A=10 \text{ AND } B=4 \text{ AND } C=6}(R)$

Exercise 2 Ans

$T(R) = 10000$, $B(R) = 400$, $V(R,A) = 500$, $V(R,B)=300$, $V(R,C)=100$. R has a clustering index on A and non-clustering index on B .

Calculate the cost in terms of a) the number of tuples expected in the result and b) the IOs it will take to retrieve the result in the following cases (assume index to be in memory):

1. $\sigma_{A<10}(R)$

- a. Tuples: $T(R)/3 = 10000/3$
- b. IOs: $B(R)/3 = 400/3$

2. $\sigma_{B=20}(R)$

- a. Tuples: $T(R)/V(R,B) = 10000/300$
- b. IOs: $T(R)/V(R,B) = 10000/300$

3. $\sigma_{C=40}(R)$

- a. Tuples: $T(R)/V(R,C) = 10000/100$
- b. IOs: $B(R) = 400$

What would be a cost effective plan to execute this query:

4. $\sigma_{A=10 \text{ AND } B=4 \text{ AND } C=6}(R)$

4.

All the tuples satisfying $A=10$ ($T(R)/500$) can be brought to memory in $B(R)/V(R,A) = 400/500$ IOs, ie in one IO. Therefore its better to use the index on A to get these tuples satisfying $A=10$ to memory. Once in memory we can look for the ones among these which satisfy the other two conditions, ie $B=4$ and $C=6$.

Exercise 3

This question is created from the DSCB example mentioned earlier.

Suppose $B(R) = 1000$, and $T(R) = 20,000$.

R is clustered and A is one of the attributes of R , and there is a non-clustering-index on A . What would be a good way to execute the query $\sigma_{A=0}(R)$.

1. Index probe costs 2 IO per probe
2. Index is in memory

Exercise 3 Ans

This question is created from the DSCB example mentioned earlier.

Suppose $B(R) = 1000$, and $T(R) = 20,000$.

R is clustered and A and B are the attributes of R with $V(R,A) = 10$, and there is a non-clustering-index on A. What would be a good way to execute the query $\sigma_{A=0}(R)$.

There can be two ways of executing this query, simply reading the whole Relation ie $B(R)$ IOs or using the index.

1. Index probe costs 2 IO per probe
 - a. Reading the relation = $B(R) = 1000$ IOs
 - b. Using Index : $2 + T(R)/V(R,A) = 2 + 2000 = 2002$ IOs
2. Index is in memory
 - a. Reading the relation = $B(R) = 1000$ IOs
 - b. Using Index : $T(R)/V(R,A) = 2000$ IOs

Please repeat this question for attribute B, with $V(R,B) = 100$ and a clustering index on B and query: $\sigma_{B<10}(R)$.

⇒ **Cheaper to read the whole relation than to use index in both scenarios.**

Exercise 4

Relation R : $B(R) = 10,000$

Relation S : $B(S) = 4000$

Given that R and S are sorted and contain no duplicates for join attribute, what Join would you suggest. What would be the

Exercise 4 Ans

Relation R : $B(R) = 10,000$

Relation S : $B(S) = 4000$

Given that R and S are sorted and contain no duplicates for join attribute, what Join would you suggest. What would be the minimum memory requirement for such a join.

Since the question specifically states there are no duplicates we can do the join with even two buffers using merge join. (The number of duplicates can place a constraint on the number of buffers required for merge join).

Exercise 4

Estimate the number of tuples for:

$R(a, b, c) \bowtie_{R.b=S.d \text{ AND } R.c=S.e} S(d, e, f)$

with

$R(a, b, c)$	$S(d, e, f)$
$T(R) = 1000$	$T(S) = 2000$
$V(R, b) = 20$	$V(S, d) = 50$
$V(R, c) = 100$	$V(S, e) = 50$

We can think of this join as a natural join if we regard R.b and S.d as the same attribute and also regard R.c and S.e as the same attribute. Then using the same rules we follow for a natural join:

$$\begin{aligned}\text{No. of Tuples} &= \frac{T(R)*T(S)}{\max\{V(R,b), V(S,d)\} * \max\{V(R,c), V(S,e)\}} \\ &= \frac{1000*2000}{(50*100)} = 400\end{aligned}$$

Note: Size of each resultant tuple = $S(R) + S(S)$

Exercise 5

Consider the join $R(a, b, c) \bowtie S(b, c, d) \bowtie U(b, e)$,

$R(a, b, c)$	$S(b, c, d)$	$U(b, e)$
$T(R) = 1000$	$T(S) = 2000$	$T(U) = 5000$
$V(R, a) = 100$		
$V(R, b) = 20$	$V(S, b) = 50$	$V(U, b) = 200$
$V(R, c) = 200$	$V(S, c) = 100$	
	$V(S, d) = 400$	
		$V(U, e) = 500$

And the following sizes:

$R.a = 10$ bytes

$R.b = S.b = U.b = 5$ bytes

$R.c = S.c = 20$ bytes

$S.d = 50$ bytes

$U.e = 22$ bytes

What will be the size of the join result?

Exercise 5 Ans

1. The number of tuples =
$$\frac{\left(\frac{1000 * 2000}{50 * 200} \right) * 5000}{200}$$
$$= 5000$$
2. Size of each tuple = $S(R) + S(S) + S(U) - S(R.b) - S(R.c) - S(U.b)$
$$= 35 + 75 + 27 - 5 - 20 - 5$$
$$= 107 \text{ bytes}$$

Therefore size of result = **5000x107** bytes.