

## Running Example

Two Relations R and S

$T(R) = 60000$

$T(S) = 2000$

Block size = 4000 bytes

$S(R) = 100$  bytes

$S(S) = 500$  bytes

Main Memory = 126 Buffers

---

### Example 1a: Join Selection for Large Relations

Two Relations R and S:

- $T(R)=80,000$ ,  $T(S)=3,000$
- Block size = 4000 bytes,  $S(R)=120$ ,  $S(S)=600$  bytes
- Main Memory = 150 Buffers

Given the above details, what join algorithm would you suggest? Calculate the total number of I/O operations required.

---

### Example 1b: Join Selection for Medium Relations

Two Relations R and S:

- $T(R)=30,000$ ,  $T(S)=1,500$ , Block size = 4000 bytes,  $S(R)=80$  bytes
- $S(S)=400$  bytes, Main Memory = 100 Buffers

Given the above details, what join algorithm would you suggest? Calculate the total number of I/O operations required.

---

### Example 1c: Join Selection for Small Relations

Two Relations R and S:

- $T(R)=5,000$ ,  $T(S)=500$ , Block size = 4000 bytes,  $S(R)=100$  bytes,  $S(S)=500$  bytes
- Main Memory = 50 Buffers

Given the above details, what join algorithm would you suggest? Calculate the total number of I/O operations required.

---

### Explanation of Concepts in the Questions

- **Tuple Count  $T(R)$  and  $T(S)$ :** The number of tuples in each relation.
- **Block Size:** The size of each block in bytes.
- **Tuple Size ( $S(R)$  and  $S(S)$ ):** The size of each tuple in bytes.
- **Main Memory Buffers:** The number of buffers available in main memory for processing.

These questions test understanding of:

- **Join Algorithm Selection:** Choosing the most efficient join algorithm (e.g., nested loop join, sort-merge join, hash join) based on the given parameters.
  - **I/O Operations:** Calculating the total number of disk I/O operations required for the join.
- 

## Memory Requirement for Hash Join

Let's assume in our Hash join example from the last slide, **the number of main memory buffers available are only 6**. The rest of the specifications are the same. Now only 5 buffers are available for bucketizing.

Number of phases or passes :  $\min(B(R), B(S)) \leq M^n$ . ie  $250 \leq 6^n$ . This gives  $n = 4$ , ie we will need four passes to perform this join. The first three will be bucketizing and the last will be the join.

1. Phase 1 : Bucketizing
  - a. R : Divide the 1500 blocks into 5 buckets, reading one block at a time, we get 300 blocks per bucket.
  - b. S : Divide 250 blocks into 5 buckets, we get 50 blocks per bucket.
  - c. Cost =  $2(B(R) + B(S)) = 3500$
2. Phase 2 : re-Bucketizing
  - a. R: Further divide each of the 5 buckets into another 5, size of new buckets =  $300/5 = 60$  blocks per new bucket.
  - b. S : Further divide each of the 5 buckets into another 5, size of new buckets =  $50/5 = 10$  blocks per new bucket.
  - c. Cost =  $2(B(R) + B(S)) = 3500$
3. Phase 3 : re-re-Bucketizing
  - a. R: new bucket size =  $60/5 = 12$  blocks / bucket.
  - b. S: new bucket size =  $10/5 = 2$  blocks / bucket. **Finally we have a bucket small enough to fit in main memory with a block to spare!**
  - c. Cost =  $2(B(R) + B(S)) = 3500$
4. Phase 4 : Join
  - a. **Only Ss buckets can be kept in the memory** while the corresponding R Bucket is read one block at a time and joined with the whole of the S bucket.
  - b. Cost =  $B(R) + B(S) = 1750$

Total cost =  $3 \times 2 \times (B(R) + B(S)) + (B(R) + B(S)) = 3 \times 3500 + 1750 = 12,250$  IOs.

---

### Question 1a: Hash Join with Limited Memory

Let's assume in a Hash Join example, the number of main memory buffers available is only **8**. The rest of the specifications are as follows:  $B(R)=2000$  blocks,  $B(S)=500$  blocks, Number of buffers for bucketizing: **7**

- a) Calculate the number of phases or passes required for the Hash Join.
  - b) Describe the bucketizing process for each phase and calculate the I/O cost for each phase.
  - c) What is the total I/O cost for the entire Hash Join operation?
- 

### Question 1b: Hash Join with Smaller Relations

Let's assume in a Hash Join example, the number of main memory buffers available is only **5**. The rest of the specifications are as follows:  $B(R)=1000$  blocks,  $B(S)=200$  blocks, Number of buffers for bucketizing: **4**

- a) Calculate the number of phases or passes required for the Hash Join.
- b) Describe the bucketizing process for each phase and calculate the I/O cost for each phase.
- c) What is the total I/O cost for the entire Hash Join operation?

1b  
 $M = 5$ , buffer = 9.  
 $B(R) = 1000 \text{ blocks}$ ,  $B(S) = 200 \text{ blocks}$   
 $\min(1000, 200) \leq 5^n \rightarrow 200 \leq 5^n \rightarrow 5^4 = 625$   
 phases = 11  
 a) R:  $\frac{200}{5} = 40 \text{ blocks/bucket}$ .  
 b) S:  $\frac{1000}{5} = 200 \text{ blocks/bucket}$ .  
 c)  $C = 2(1200) = 2400 \text{ I/O}$   
 2) a) R:  $\lceil \frac{200}{5} \rceil = 40 \text{ blocks/bucket}$ .  
 b) S:  $\lceil \frac{1000}{5} \rceil = 200 \text{ blocks/bucket}$ .  
 c) cost:  $2400 \text{ I/O}$   
 3) a) R:  $\lceil \frac{40}{5} \rceil = 8 \text{ blocks/bucket}$ .  
 b) S:  $\lceil \frac{200}{5} \rceil = 40 \text{ blocks/bucket}$ .  
 $\text{cost} = 2400 \text{ I/O} \approx 480 \text{ fit in memory}$   
 4) Join:  $1000 + 200 = 1200 \text{ I/O}$   
 Total  $\Rightarrow 3(2400) + 1200 = 7200 + 1200 = \frac{8400}{10}$

---

### Question 1c: Hash Join with Larger Relations

Let's assume in a Hash Join example, the number of main memory buffers available is only **10**. The rest of the specifications are as follows:  $B(R)=3000$  blocks,  $B(S)=1000$  blocks, Number of buffers for bucketizing: **9**

- a) Calculate the number of phases or passes required for the Hash Join.
  - b) Describe the bucketizing process for each phase and calculate the I/O cost for each phase.
  - c) What is the total I/O cost for the entire Hash Join operation?
- 

### Explanation of Concepts in the Questions

1. **Hash Join:** A join algorithm that uses hashing to partition relations into smaller buckets that fit in memory.
2. **Phases or Passes:** The number of times data must be read/written to disk during bucketizing and joining.
3. **Bucketizing:** Dividing the relations into smaller buckets using a hash function.
4. **I/O Cost:** The total number of disk I/O operations required for the join, calculated as:
  - o **Bucketizing Cost:**  $2 \times (B(R) + B(S))$  per phase.
  - o **Join Cost:**  $B(R) + B(S)$  for the final join phase.

# Exercise 1

Relation R :  $B(R) = 10,000$

Relation S :  $B(S) = 4000$

Calculate the number of phases and the IOs for a Hash Join if R and S are contiguous for:

- a)  $M = 100$
  - b)  $M = 50$
  - c)  $M = 30$
- 

## Exercise 1a

Relation R:  $B(R)=12,000$ .      Relation S:  $B(S)=5,000$

Calculate the **number of phases** and the **I/O cost** for a **Hash Join** if **R and S are contiguous** for:

- a)  $M=120$ , b)  $M=60$ , c)  $M=40$
- 

## Exercise 1b

Relation R:  $B(R)=15,000$  Relation S:  $B(S)=6,000$

Compute the **number of phases** and the **I/O cost** for a **Hash Join** if **R and S are contiguous** for:

- a)  $M=150$ , b)  $M=80$ , c)  $M=50$
- 

## Exercise 1c

Relation R:  $B(R)=8,000$  Relation S:  $B(S)=3,500$

Determine the **number of phases** and the **I/O cost** for a **Hash Join** if **R and S are contiguous** for:

- a)  $M=90$
  - b)  $M=45$
  - c)  $M=25$
-

## Exercise 2

Relation R :  $B(R) = 10,000$

Relation S :  $B(S) = 4000$

Calculate the number of phases and the IOs for a Merge Join if R and S are contiguous. Optimize whenever possible.

- a)  $M = 101$
- b)  $M = 50$
- c)  $M = 30$

### Exercise 2a

Relation R:  $B(R)=12,000$  Relation S:  $B(S)=5,000$

Calculate the **number of phases** and the **I/O cost** for a **Merge Join** if **R and S are contiguous**. Optimize whenever possible.

- a)  $M=120$
  - b)  $M=60$
  - c)  $M=40$
- 

### Exercise 2b

Relation R:  $B(R)=15,000$  Relation S:  $B(S)=6,000$

Calculate the **number of phases** and the **I/O cost** for a **Merge Join** if **R and S are contiguous**. Optimize whenever possible.

- a)  $M=150$
  - b)  $M=80$
  - c)  $M=50$
- 

### Exercise 2c

Relation R:  $B(R)=8,000$  Relation S:  $B(S)=3,500$

Calculate the **number of phases** and the **I/O cost** for a **Merge Join** if **R and S are contiguous**. Optimize whenever possible.

- a)  $M=90$
- b)  $M=45$
- c)  $M=25$

## Exercise 3

Relation R : B(R) = 10,000, T(R) = 100,000

Relation S : B(S) = 4000, T(S) = 32,000

M = 100. Calculate IOs:

- a) for Hash Join, if the relations are not contiguous
- b) for Hash Join, if the relations are sorted
- c) for Merge Join, if the relations are sorted
- d) And Minimum memory requirement for One Pass Join if the relations are Contiguous.

### Exercise 3a

Relation R: B(R)=12,000, T(R)=120,000. Relation. S: B(S)=5,000, T(S)=40,000 Available **main memory blocks**: M=120

Calculate the I/O costs for:

- a) Hash Join when the relations are **not contiguous**.
  - b) Hash Join when the relations are **sorted**.
  - c) Merge Join when the relations are **sorted**.
  - d) The **minimum memory requirement** for **One-Pass Join** when the relations are **contiguous**.
- 

### Exercise 3b

Relation R: B(R)=15,000, T(R)=150,000. Relation S: B(S)=6,000, T(S)=48,000. Available **main memory blocks**: M=150

Calculate the I/O costs for:

- a) Hash Join if the relations are **not contiguous**.
  - b) Hash Join if the relations are **already sorted**.
  - c) Merge Join if both relations are **sorted**.
  - d) The **minimum memory requirement** for a **One-Pass Join**, assuming **contiguous storage**.
- 

### Exercise 3c

Relation R: B(R)=8,000, T(R)=80,000. Relation. S: B(S)=3,500, T(S)=28,000. Available **main memory blocks**: M=80

Compute the I/O costs for:

- a) Hash Join if the relations are **not stored contiguously**.
  - b) Hash Join if both relations are **sorted**.
  - c) Merge Join assuming the relations are **sorted**.
  - d) The **minimum memory requirement** for a **One-Pass Join** when the relations are **contiguous**.
-

## Joins: Iteration vs. Hash vs. Merge

Joins are a fundamental operation in relational databases, used to combine records from two or more tables based on a common key. There are multiple ways to implement joins efficiently, depending on the **size of the tables, availability of indexes, and amount of memory available**.

The three major types of join algorithms are:

1. **Nested Loop Join (Iteration-based Join)**
2. **Hash Join**
3. **Merge Join (Sort-Merge Join)**

---

### 1. Nested Loop Join (Iteration-based Join)

#### Concept:

This is the most basic way to perform a join. It involves iterating through one relation (outer loop) and, for each row, scanning the entire second relation (inner loop) to find matching tuples.

#### Types:

- **Brute Force Nested Loop:** No indexing, scanning every combination of tuples.
- **Indexed Nested Loop:** Uses an index on the inner relation to speed up lookups.

#### Time Complexity:

- **Without Index:**  $O(N \times M)$  (where  $N$  and  $M$  are the sizes of the two relations)
- **With Index:**  $O(N \times \log M)$  (if an index exists on the inner relation)

#### Example:

Consider two tables:

```
sql
SELECT *
FROM Employees E, Departments D
WHERE E.DepartmentID = D.DepartmentID;
```

- **Brute Force:** For each employee, scan all departments.
- **With Index:** If `DepartmentID` in `Departments` is indexed, lookup is much faster.

#### I/O Cost:

For relations  $R$  (outer) and  $S$  (inner), assuming:

- $B(R) = 10,000$  blocks
- $B(S) = 4,000$  blocks
- $\text{Memory} = M = 100$  blocks



### Brute Force Nested Loop I/O Cost:

$$B(R) \times B(S) = 10,000 \times 4,000 = 40,000,000 \text{ I/Os}$$

### Indexed Nested Loop I/O Cost (if index exists on S):

$$B(R) + (T(R) \times \text{Index Cost})$$

### Pros & Cons:

- ✓ Simple to implement
  - ✓ Works without sorting or hashing
  - ✗ Very slow for large datasets
  - ✗ Not efficient unless indexes exist
- 

## 2. Hash Join

### Concept:

Uses **hashing** to divide records into buckets and then joins corresponding buckets.

### Phases:

#### 1. Partition Phase:

- Use a hash function  $h(x)$  to distribute both relations into **buckets** stored in memory/disk.

#### 2. Join Phase:

- Match records in the same bucket.

### Time Complexity:

$$O(N + M)$$

(Better than Nested Loop!)

### Example:

sql

↪ Kopieren ⚙ Bearbeiten

```
SELECT *
FROM Orders O, Customers C
WHERE O.CustomerID = C.CustomerID;
```

- A **hash function** is applied to `CustomerID` to partition both `Orders` and `Customers` into buckets.
- Matching is only done **within buckets**, reducing comparisons.

## I/O Cost:

For R (10,000 blocks) and S (4,000 blocks):

- Partitioning:  $B(R) + B(S) = 10,000 + 4,000 = 14,000$
- Probing:  $B(R) + B(S) = 14,000$

Total I/Os = 28,000

## Pros & Cons:

- Faster than Nested Loop, especially for large datasets
- Does not require sorting
- Needs more memory
- Performance drops if hash function creates skewed buckets

---

## 3. Merge Join (Sort-Merge Join)

### Concept:

Sorts both relations on the join key and then merges them in a single pass.

### Phases:

#### 1. Sorting Phase:

- Sort both relations on the join key.

#### 2. Merge Phase:

- Scan both sorted relations and match tuples.

### Time Complexity:

$$O(N \log N + M \log M)$$

(because of sorting)

$$+ O(N + M)$$

(for merging)

### Example:

sql

 Kopieren  Bearbeiten

```
SELECT *
FROM Students S, Enrollments E
WHERE S.StudentID = E.StudentID;
```

- Both Students and Enrollments are sorted on StudentID.
- They are merged efficiently by scanning both tables in order.

### I/O Cost:

For R (10,000 blocks) and S (4,000 blocks):

- Sorting:  $2B(R) \log_M B(R) + 2B(S) \log_M B(S)$
- Merging:  $B(R) + B(S) = 14,000$

Total I/Os  $\approx 30,000$  (depends on sorting efficiency)

### Pros & Cons:

- Best when relations are already sorted
  - Efficient for large datasets
  - Works well for range queries
  - Expensive if sorting is required
  - Needs enough memory to hold sorted runs
- 

## Comparison Table

Join Type	Best Use Case	Time Complexity	Memory Required	I/O Cost
Nested Loop	Small tables or indexed joins	$O(N \times M)$	Low	Very high
Hash Join	Large tables, equality joins	$O(N + M)$	Medium to High	Medium
Merge Join	Already sorted data, range joins	$O(N \log N + M \log M)$	Medium	Medium

---

## When to Use Each Join Type?

- Use Nested Loop if relations are small or indexes exist.
  - Use Hash Join for large, unsorted tables when equality conditions are used.
  - Use Merge Join when both relations are sorted (or sorting is acceptable).
- 

## Conclusion

Joins are a critical operation in query processing, and choosing the right algorithm is crucial for performance.

- Nested Loop is simple but inefficient.
- Hash Join is fast for large, unsorted data.
- Merge Join is great for sorted relations.